

Training a Transformer from Scratch

A modern decoder-only GPT (RoPE, RMSNorm, SwiGLU) and a small scaling study

Marcos Lahoz

May 29, 2026

Abstract

We implement a decoder-only Transformer language model from scratch in PyTorch — using only tensor primitives, no pretrained or library model classes — with the components found in current LLMs: pre-norm **RMSNorm**, **rotary position embeddings** (RoPE), **SwiGLU** feed-forward blocks, weight tying and a **KV cache** for generation. We verify the implementation on a synthetic copy task (98.4% exact-copy accuracy), train it as a word-level language model on TinyStories, and run a small **scaling study** relating validation loss to model size. The code is a compact, reproducible reference for how a modern GPT is built and trained.

1 Architecture

The model is a stack of L pre-norm Transformer blocks. For hidden states x , each block computes

$$x \leftarrow x + \text{Attn}(\text{RMSNorm}(x)), \quad x \leftarrow x + \text{SwiGLU}(\text{RMSNorm}(x)).$$

RMSNorm [3] normalises by the root-mean-square of the activations (no mean subtraction, no bias). **RoPE** [2] rotates the query/key vectors by position-dependent angles, injecting relative position directly into the attention dot-product. Attention is the standard causal multi-head self-attention; for generation we keep a per-layer **KV cache** so each new token costs $O(T)$ rather than $O(T^2)$. The MLP is **SwiGLU** [4], $\text{SwiGLU}(x) = W_{\text{down}}(\text{SiLU}(W_{\text{gate}}x) \odot W_{\text{up}}x)$. The token-embedding and output projection share weights (weight tying).

2 Implementation validation

Before training on text we validate the full pipeline (forward, loss, backward, and cached generation) on a *copy task*: the input is a random length-8 sequence, a separator, and the same sequence again; only the copied part is scored. Solving it requires attending back to the prefix. A 2-layer, 64-dim model (107k parameters) reaches a loss of 2×10^{-4} in 300 steps and **98.4%** exact-copy accuracy when generating with the KV cache — confirming attention, RoPE, masking and sampling are correct. Reproduce with `python scripts/smoke_test.py`.

3 Language modeling

We BPE-tokenise (GPT-2 vocabulary) a TinyStories stream into a flat token array and sample fixed-length blocks. Training uses AdamW with decoupled weight decay, a cosine learning-rate schedule with warmup, gradient clipping, gradient accumulation and mixed precision (bf16/fp16). Validation loss and perplexity are logged periodically; `scripts/generate.py` samples text from a checkpoint. (Fill final perplexity and qualitative samples from your run.)

4 Scaling study

We train four sizes (tiny/small/medium/large, from ~ 0.5 M to ~ 85 M non-embedding parameters) for a fixed iteration budget and record the final validation loss. We expect the familiar monotone, roughly power-law decrease of loss with model size.

[Figure *scaling.png* is produced by *scripts/scaling_study.py*.]

size	params (M)	val loss	val ppl
tiny	...	...	...
small	...	...	...
medium	...	...	...
large	...	...	...

Table 1: Scaling summary (produced as `scaling.csv`).

5 Discussion

The project is a clean, dependency-light reference implementation of a modern GPT that is both *correct* (validated end-to-end) and *scalable*. Natural extensions: grouped-query attention, FlashAttention / `scaled_dot_product_attention`, μ P or careful LR scaling, longer context, and fitting an explicit scaling law $L(N) = L_\infty + (N_0/N)^\alpha$ to the measured points.

References

- [1] Vaswani et al. (2017). *Attention Is All You Need*. NeurIPS.
- [2] Su et al. (2021). *RoFormer: Enhanced Transformer with Rotary Position Embedding*.
- [3] Zhang & Sennrich (2019). *Root Mean Square Layer Normalization*. NeurIPS.
- [4] Shazeer (2020). *GLU Variants Improve Transformer*.
- [5] Kaplan et al. (2020). *Scaling Laws for Neural Language Models*.